

# Time Travelling PostgreSQL

PGDay/MED 2024-12-05  
Stefanie Janine Stölting





# About Me



Owner of Tallinn/Estonia based ProOpenSource OÜ

Website: [www.ProOpenSource.eu](http://www.ProOpenSource.eu)

Blog: [www.ProOpenSource.it](http://www.ProOpenSource.it)

Mastodon: <https://digitalcourage.social/@sjstoelting>


Telegram: <https://t.me/postgreschat>



# About Me



Member of the PGEU Diversity Committee



PostgreSQL Europe Diversity Task Force



# What is Time Travelling







# Definition



Being able to return data as it has been at a certain point in time.

That can obviously only be data saved in the past.

And we need dates and times stored with data.




Photo by [Tim Green](#)



# History



PostgreSQL supported time travel as it was implemented in its core up to version 6.2. It has been removed in [PostgreSQL 6.3](#).





# Other Ideas



Hans-Jürgen Schöning from CYBERTEC has written about implementing time travelling in his blog post in 2019 about [AS OF-queries](#).

I have tested them, but they seem not to work with PostgreSQL 17.

The creation of the trigger fails.



# Other Ideas



The problem here is:

```
CREATE FUNCTION version_trigger() RETURNS trigger AS
$
...
...
END;
$ LANGUAGE plpgsql;
```

The solution with PostgreSQL 17 is:

```
CREATE FUNCTION version_trigger()
  RETURNS TRIGGER
  LANGUAGE PLPGSQL
AS
$$
...
...
END;
$$
```





# What is covered?



Table partitioning by range

Primary key with partitioned tables

Updating records with the possibility to access the previous records as they have been at a certain point in time

Deleting records with the possibility to access the previous records as they have been at a certain point in time



# Let's start



Photo by [Wim van 't Einde](#) on [Unsplash](#)



# Primary Key - Partitioning



PostgreSQL does not support unique indexes or primary keys over partitioned tables.

But there is a solution to solve the primary key problem without any extensions in PostgreSQL.





# Primary Key - Table



The solution is somewhat easy.

In addition to the partitioned table we also need a table, that will hold the primary key.

The partitioned table needs a foreign key referencing the primary key table.



# Preparation



I stated we do not need an extension.

Well, we need one to get speedy access to the data: `btree_gist`.

`btree_gist` is an index that is part of the contrib package and available on all supported platforms.

```
CREATE EXTENSION IF NOT EXISTS btree_gist;
```





# Time Travel Tables



```
CREATE TABLE timetravel_pk (  
    timetravelid BIGINT GENERATED ALWAYS AS IDENTITY,  
    created timestamp with time zone NOT NULL,  
    CONSTRAINT timetravle_pk_pk PRIMARY KEY (timetravelid)  
);
```

```
CREATE TABLE timetravel (  
    timetravelid BIGINT NOT NULL,  
    changed TSTZRANGE NOT NULL DEFAULT tstzrange(clock_timestamp(), 'INFINITY', '[]'),  
    data_text TEXT,  
    data_json JSONB,  
    deleted BOOLEAN NOT NULL DEFAULT FALSE,  
    CONSTRAINT timetravelid_fk FOREIGN KEY (timetravelid)  
        REFERENCES timetravel_pk(timetravelid)  
    ) PARTITION BY RANGE (lower(changed));
```



# Indexes



```
CREATE INDEX timetravel_changed_idx  
  ON timetravel  
  USING gist  
  (changed)  
;
```

```
CREATE INDEX timetravel_timetravelid_idx  
  ON timetravel  
  USING btree  
  (timetravelid)  
;
```


```
CREATE INDEX timetravel_not_deleted_idx  
  ON timetravel  
  USING btree  
  (deleted)  
  WHERE NOT deleted  
;
```



# Partition Management Table



```
CREATE TABLE timetravel_part_vals (  
  part_year SMALLINT NOT NULL,  
  start_value TIMESTAMP WITH TIME ZONE NOT NULL,  
  end_value TIMESTAMP WITH TIME ZONE NOT NULL,  
  CONSTRAINT timetravel_part_vals_pk PRIMARY KEY (part_year)  
);
```



# Partition Management Table



This table is needed to handle partition information. It will later be used to check, if a partition has to be created.

A function will use that table to create partitions.

The function will be executed manually.

In real life I would use a cron job, for example with `pg_cron`.



# Show Time



Showing slides with SQL is a bit boring.








# Link List



Slide and sources are available on GitLab and GitHub:

<https://gitlab.com/sjstoelting/talks/>

<https://github.com/sjstoelting/talks/>



# Time Traveling PostgreSQL



This document by [Stefanie Janine Stölting](#) is covered by the [Creative Commons Attribution 4.0 International](#)